


Continuous-Time Generative Modelling & Chemistry

Chenghao Liu
Caltech



LASSLGRRIQLRGINKLTDNLTSEKEGSGPPNWELRNKAQGDIIHNEKENAEAHKDL DGRDRFGPRDWNKDKVPEG
GTVVNITEGAGKAFYRPEGYTDWKQLILDWCCNVRGASGDLNGSHGLSAGPFWVVLKIVSCNVKERKRKVGVAVE



Modern Image & Video Generation




Wan




Beyond Images and Text

Scientific Data




SE(3) invariant
Protein structure generation

Robotics




SO(2) invariant
Block stacking

Information Geometry



Fisher-Rao geometry
On the probability Simplex

Climate Modeling



Spherical Geometry $\mathbb{S}^2$

Outline

Part I. Simulation(-free) Generative Models


Part II. Fine-tuning & Steering

Part III. Applications to Chemistry




Problem Setting: Generative Modeling

- **Unknown:** data distribution q




Problem Setting: Generative Modeling

- **Unknown:** data distribution q
- **Given:** samples $x_1 \sim q$




Problem Setting: Generative Modeling

- **Unknown:** data distribution q
- **Given:** samples $x_1 \sim q$



Problem Setting: Generative Modeling



- **Unknown:** data distribution q
- **Given:** samples $x_1 \sim q$



Goal: learn *a sampler* from the unknown q

Problem Setting: Generative Modeling

- **Unknown:** data distribution q
- **Given:** samples $x_1 \sim q$
- **Learn:** neural network with parameters θ




Goal: find parameters θ s.t. $p_\theta \approx q$

Deep Generative Modelling


Easy to sample from

$$x_0 \sim p$$



Generator
 ψ_θ

$$x_1 \sim q$$




Sampling

$$x_0 \sim p$$
$$\psi_\theta(x_0) \sim q$$

How to model ψ_θ ?

Density Estimation

$$p_\theta \approx q$$



Deep Generative Modelling

Fidelity

Samples should look like they came from q

Coverage

Model all of q , not just its easy modes, and ideally give $\log p(x)$

Speed

Train fast (& stable). Call a few times during inference.

Deep Generative Modelling

1

One-step maps

GAN, VAE,
Normalizing flows:
one big jump from
noise to data.

2

Neural ODEs

ODEs: infinitely many
infinitesimal steps.

3

Diffusion

Add noise gradually;
learn to undo it. It
works — at a price.

4

Flow matching

Choose the path;
Regress its velocity.
No simulation.

5

One step, again

Learn the journey at
once: consistency &
flow maps.

Summary of Modern Deep Generative Models

Generative Model

$$(\psi_{\theta}, p_{\theta})$$


Autoregressive Models

$$p_{\theta}(x) = \prod_{i=1}^d p_{\theta}(x_i | x_{<i})$$

■


- Likelihood-based
- Sequential generation — scales with dimension
- SOTA in language modeling

GANs



- Fast Sampling!
- No Likelihood
- Delicate training

VAEs



- Fast Sampling!
- Optimize ELBO
- No Likelihood
- Inferior performance

Why are diffusion and flow useful?

GANs are still competitive performance-wise even today

Class-Conditional ImageNet 512×512					
Model	FID↓	sFID↓	IS↑	Precision↑	Recall↑
BigGAN-deep [10]	8.43	8.13	177.90	0.88	0.29
StyleGAN-XL [57]	2.41	4.06	267.75	0.77	0.52
Mask-GIT [12]	7.32	-	156.0	-	-
ADM [19]	23.24	10.19	58.06	0.73	0.60
ADM-G, ADM-U	3.85	5.86	221.72	0.84	0.53
Simple Diffusion(U-Net) [28]	4.28	-	171.0	-	-
Simple Diffusion(U-ViT, L)	4.53	-	205.3	-	-
VDM++ [36]	2.65	-	278.1	-	-
DiT-XL(cfg = 1.5) [50]	3.04	5.02	240.82	0.84	0.54
SiT-XL(cfg = 1.5, SDE)	2.62	4.18	252.21	0.84	0.57

It's about training stability and inference flexibility

Normalizing Flows

Sample from some complicated distribution by sampling from a simple distribution then applying U

- Begin with a simple distribution


$$p_{t_0}(x) \sim \mathcal{N}(0, 1)$$

- Apply an *invertible* transformation(s)



$$x_{t_1} = U(x_{t_0})$$

- Use change of variables to calculate probability

$$\log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \log \det \left| \frac{\partial U}{\partial x_{t_0}} \right|$$



Normalizing Flows



Normalizing Flows are Capable Generative Models

Shuangfei Zhai¹ Ruixiang Zhang¹ Preetum Nakkiran¹ David Berthelot¹ Jiatao Gu¹ Huangjie Zheng¹
Tianrong Chen¹ Miguel Angel Bautista¹ Navdeep Jaitly¹ Josh Susskind¹

Deep Normalizing Flows

- ✓ exact log-likelihood
- ✓ sampling in one pass
- ✗ architecture restricted to invertible blocks
- ✗ latent dimension must equal data dimension

Apply a series of transformations

$$x_{t_1} = U(x_{t_0}) \longrightarrow x_{t_N} = u_N \circ u_{N-1} \circ \cdots \circ u_1(x_{t_0})$$

Use change of variables to calculate probability

$$\log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \log \det \left| \frac{\partial U}{\partial x_{t_0}} \right| \longrightarrow \log p_{t_N}(x_{t_N}) = \log p_{t_0}(x_{t_0}) - \sum_{n=1}^N \log \det \left| \frac{\partial u_n}{\partial x_{t_{n-1}}} \right|$$

Continuous Normalizing Flows

**Infinitely many layers
= an ordinary differential equation**

Apply a series of transformations


$$x_{t_N} = u_N \circ u_{N-1} \circ \cdots \circ u_1(x_{t_0}) \longrightarrow x_{t_1} = U(x_{t_0}) = \int_{t_0}^{t_1} u(x(t), t) dt$$

Use change of variables to calculate probability

$$\log p_{t_N}(x_{t_N}) = \log p_{t_0}(x_{t_0}) - \sum_{n=1}^N \log \det \left| \frac{\partial u_n}{\partial x_{t_{n-1}}} \right| \longrightarrow \log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial u}{\partial x(t)} \right) dt$$

Constructs an invertible
transformation ψ for any f !

The catch: training requires simulation



- To compute one training likelihood, you must solve the ODE for that sample — typically ~ 100 network evaluations...
- ...while estimating the divergence along the way (Hutchinson trace estimator)...
- ...and then backpropagate through the entire solve (or solve an adjoint ODE backwards).
- **Hundreds of network calls per gradient step.**
- **Can we not solve ODEs?**

Short symmetry


Discrete flows pay in architecture (constrained layers, cheap training).

Continuous flows pay in training time (free architecture, costly simulation).


How does this look like?




Target




Density




Samples






Vector Field



Generative Models as Dynamic Systems



Dynamical Systems as Generative Models



$$\psi_t: [0,1] \times \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Time-Dependent Generator

Sampling = Simulating

$$x_0 \sim p \quad \longrightarrow \quad \text{simulate}(x_0, t) = \psi_t(x_0)$$

Deterministic and Stochastic Dynamics

Flows

ODE

$$dx_t = u_t(x_t)dt$$


Velocity field

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$


Drift


Diffusion
Coefficient


Brownian
Motion


Deterministic and Stochastic Dynamics

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field



Deterministic

$$x_t = x_0 + \int_0^t u_s(x_s)ds$$

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion


Deterministic and Stochastic Dynamics

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field



Deterministic

$$x_t = x_0 + \int_0^t u_s(x_s)ds$$

Diffusion


SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion



Stochastic

$$x_t = x_0 + \int_0^t f_s(x_s)ds + \int_0^t g_s dw_s$$

Where are my probabilities?

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = -\text{div}(p_t u_t)$$

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Where are my probabilities?

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = -\text{div}(p_t u_t)$$

Yes!

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing...

Can we build a generative model with these?

Diffusion and score models

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion Coefficient

Brownian Motion

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing...

Diffusion and score models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data $\rightarrow$ Noise

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing...

Diffusion and score models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data $\rightarrow$ Noise

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t) dt + g_t d\bar{w}_t$$

Noise $\rightarrow$ Data

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing... The Score!

Diffusion and score models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_tdw_t$$

Data $\rightarrow$ Noise

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t) dt + g_t d\bar{w}_t$$

Noise $\rightarrow$ Data

Learn the score by regressing to conditional scores:

$$\min_{\theta} \mathbb{E}_{p_{data}, p_t(x|x_{data})} [\| s_t^{\theta}(x) - \nabla \log p_t(x|x_{data}) \|^2]$$

Simulation-free

Known SDEs: Variance Exploding

Variance Preserving

Where are my probabilities?

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = -\text{div}(p_t u_t)$$

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Fokker-Planck Equation

$$\partial_t p_t = -\text{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Learn: score $\nabla \log p_t$

- Only Gaussian source
- Solution asymptotically reaches source

Diffusion is a many-step VAE

encoder: fixed Gaussian noising $q(x_{t'}|x_t)$ — nothing to learn



decoder: learned denoising $p_\theta(x_t|x_{t'})$ — one shared network

- Latents = the noisy chain x_t . The encoder is the fixed noising process.
- The decoder undoes one small step at a time; each step is a nearly-Gaussian, nearly-linear denoising problem

Flow models

Flows

ODE

$$dx_t = u_t(x_t)dt$$


Velocity field


The Continuity Equation

$$\partial_t p_t = -\text{div}(p_t u_t)$$

Yes!

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$



$\psi_t(x)$ is smooth with smooth
inverse defined by $-u_t(x)$

Where are my probabilities?

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$

The Continuity Equation

$$\partial_t p_t = -\text{div}(p_t u_t)$$

Learn: velocity field u_t

- Universal transformation between densities
- Defined on finite time interval

Diffusion

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Liouville Equation

$$\partial_t p_t = -\text{div}(p_t(f_t - \frac{1}{2} g_t^2 \nabla \log p_t))$$

Learn: score $\nabla \log p_t$

- Only Gaussian source
- Solution asymptotically reaches source

Flow vs Flow matching models

Velocity field u_t is unknown: normally need simulation

$$\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t, x \sim p_t} \|v_\theta(x, t) - u_t(x)\|^2 \quad \log p_1(x_1) = \log p(x_0) + \int_1^0 \text{div}(u_t(x_t)) dt$$
$$x_t = x_1 + \int_1^t u_s(x_s) ds$$

Hacks: condition on a single datapoint

Per pair (x_0, x_1) , the path and its velocity are known in closed form:

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t, x_1, x \sim p_t(\cdot | x_1)} \|v_\theta(x, t) - u_t(x | x_1)\|^2$$

$$\nabla_\theta \mathcal{L}_{\text{FM}} = \nabla_\theta \mathcal{L}_{\text{CFM}}$$

Intuition: regress on single-pair velocities; the network's own averaging recovers the marginal field.

The simplest path is (nearly) the best one

Linear interpolation

Conditional path: a straight line, constant velocity

$$x_t = (1 - t) x_0 + t x_1, \quad u_t(x_t | x_0, x_1) = x_1 - x_0$$


The entire training loop

```
x0 = randn(); x1 = data(); t = rand()
```

```
xt = (1-t)*x0 + t*x1
```

```
loss = || v(xt, t) - (x1 - x0) ||2
```

Why straightness pays: Euler's error comes from curvature. Straighter marginal trajectories $\Rightarrow$ fewer, larger steps at sampling time.



Flow vs diffusion

- For Gaussian paths $x_t = \alpha_t x_1 + \sigma_t x_0$, velocity and score are linearly interconvertible.
- They are almost always the same in practice, especially during training.
- The interesting question is whether to use noise during inference.
 - Deterministic is slightly faster
 - Stochastic is slightly better performance
- Two applications where it does matter
 - Manifolds, e.g. molecules
 - Sparse priors, e.g. image-to-image


Consistency models: Fewer steps

$$f_{\theta}(x_t, t) \approx \Psi_{t \rightarrow 1}(x_t)$$

Self-consistency + boundary condition

$$f_{\theta}(x_t, t) = f_{\theta}(x_s, s), \quad f_{\theta}(x_1, 1) = x_1$$

The trick: Network returns same value for any two points in the trajectory, and at $t=1$ it returns the data. Agreement between every neighboring pair chains into agreement between all pairs -> few steps.



f_{θ} maps every point on a trajectory to its endpoint

Train by consistency distillation (from a pretrained field) or consistency training (from scratch).

Consistency Models

Yang Song¹ Prafulla Dhariwal¹ Mark Chen¹ Ilya Sutskever¹

ONE STEP DIFFUSION VIA SHORTCUT MODELS


Kevin Frans
UC Berkeley
kvfrans@berkeley.edu

Danijar Hafner
UC Berkeley

Sergey Levine
UC Berkeley

Pieter Abbeel
UC Berkeley

Mean Flows



Idea:
Condition on current and
target time!

$$u_{\theta}(x_t, t)$$

↓

$$u_{\theta}(x_t, r, t)$$

How to learn $u_\theta(x_t, r, t)$

$$(t - r)u(x_t, r, t) = \int_r^t v(x_t, \tau) d\tau$$

Derivative w.r.t. t

$$u(x_t, r, t) + (t - r) \frac{d}{dt} u(x_t, r, t) = v(x_t, t)$$

Rearrange

$$u_{target}(x_t, r, t) = v(x_t, t) - (t - r) \frac{d}{dt} u(x_t, r, t)$$

$$L(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \| u_t^\theta(x, r, t) - \text{stopgrad}(u_{target}) \|^2$$

How to learn $u_\theta(x_t, r, t)$

$$(t - r)u(x_t, r, t) = \int_r^t v(x_t)$$

Derivative w.r.t. t

$$u(x_t, r, t) + (t - r) \frac{d}{dt} u(x_t, r, t) = v(x_t)$$

Rearrange

$$u_{target}(x_t, r, t) = v(x_t, t)$$

$$L(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \| u_t^\theta(x, r, t) - \text{stopgrad}(u_{target}) \|^2$$

Algorithm 1 MeanFlow: Training.

Note: in PyTorch and JAX, `jvp` returns the function output and JVP.

```
# fn(z, r, t): function to predict u
# x: training batch
```

```
t, r = sample_t_r()
e = randn_like(x)
```

```
z = (1 - t) * x + t * e
v = e - x
```

```
u, dudt = jvp(fn, (z, r, t), (v, 0, 1))
```

```
u_tgt = v - (t - r) * dudt
error = u - stopgrad(u_tgt)
```

```
loss = metric(error)
```

Algorithm 2 MeanFlow: 1-step Sampling

```
e = randn(x_shape)
x_t = e - fn(e, r=0, t=1)
```

Summary

Model	The map	Trained by	Exact log p?	Steps	Main constraint
Normalizing flow	invertible network	exact max likelihood	yes	1	invertibility requirement
VAE	stochastic decoder	ELBO (bound)	hard	1	one giant leap
GAN	free-form network	adversarial game	no	1	instability; mode dropping
Autoregressive	factorized sampling	exact max likelihood	yes	d (per dim)	sequential generation
CNF	ODE solution	exact ML via simulation	yes	many	simulation inside training
Diffusion	SDE / PF-ODE	denoising regression (= ELBO)	hard	10–100+	Gaussian path; many steps
Flow matching	ODE solution	velocity regression (CFM)	hard	5–50	Many steps
Flow map	trajectory endpoint map	self-consistency / MF identity	hard	1–4	Scaling & stability

Take-aways

- Gotta have samples!
- Generative modelling is a coupling problem: how to best map noise onto data?
- Rich training signal & inference choices drive performance.
- Flow and diffusion are basically the same.
- Most useful due to stability & flexibility.
- Many small tricks for 'best' performance.

Unsolved challenges:


- Discrete space: Autoregressive remains the best model. How to do few step generation?
How to scale diffusion model (can they be scaled?)
- Continuous space: How to do multimodal generation well, especially few step?

Outline

Part I. Simulation(-free) Generative Models

Part II. Fine-tuning & Steering

Part III. Applications to Chemistry



Pretrained models & the desired outputs



sans fine-tuning

package
package
package
package
package

By 2015-06-03 14:51:15.87945 WhyThisPost
By 2015-06-03 14:53:13.53923 WhatThisPost
By 2015-06-03 14:54:29.25974 WhenThis

...
...
...
...
...



with fine-tuning

In the quiet corners of the world, where the hum of daily life fades into the background, there exists a space where imagination and reality intertwine. It is a place where stories are not just told but felt, where every word carries the weight of possibility. Whether through the pages of a book, the brushstrokes of a painting, or the melodies of a song, this space invites us to explore, to dream, and to connect with something greater than ourselves

Pretraining gives a model

$$p_{\theta} \approx p_{data}$$


Moving from p_θ to desired outputs

$$Z = \int p_{\theta,ref}(x) \exp\left(\frac{R(x)}{\beta}\right)$$

Reward
 $R(x)$

Reward tilted distribution:

$$p^*(x) = \frac{1}{Z} p_{\theta,ref}(x) \exp\left(\frac{R(x)}{\beta}\right)$$



3 ways to sample from the target

Offline tuning

Train on some nice examples

Online tuning

Reinforcement learning to
sample from the reward

Inference steering

No training, just adapt the
sampling procedure

All sample from the *reward tilted* distribution:

$$p^*(x) \propto p_{\theta, ref}(x) \exp\left(\frac{R(x)}{\beta}\right)$$

Offline: Supervised finetuning

$$KL(p \parallel q) = \mathbb{E}_{x \sim p} \left[\log \frac{p(x)}{q(x)} \right]$$

Objective:

$$\max_{\theta} \mathbb{E}_{x \sim \mathcal{D}} [\log p_{\theta}(x)] \Leftrightarrow \min_{\theta} KL(p_{\mathcal{D}} \parallel p_{\theta})$$



Key insight: as $N \rightarrow \infty$, Best-of- N converges to $p^*(x)$

$$\propto p_{ref}(x) \exp\left(\frac{R(x)}{\beta}\right)!$$

Broadly used, but might mode collapse!

Ouyang, Long, et al. "Training language models to follow instructions with human feedback." *NeurIPS* (2022).

Offline: Direct Preferential Optimization

Dataset of pairs

$$\mathcal{D} = \left\{ \left(x_1^{(i)}, x_2^{(i)} \right) \mid x_1^{(i)} \text{ is "better than" } x_2^{(i)} \right\}$$

Assume a specific form of the optima (Bradley-Terry), then your *reward is given by your model!*

$$R(x) = \beta \log \frac{p_\theta(x)}{p_{ref}(x)} + \log Z$$

Maximum likelihood training with $\mathcal{D}$ and end up sampling from $p^*(x)$

Which response do you prefer?
Your choice will help make ChatGPT better.

 Response 1

Correct.

Explanation: The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

Question 4: Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

 Response 2

Correct.

Explanation: The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

Question 4: Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

Hard to explore

Offline: Direct Preferential Optimization

$$L_{DPO}(p_{\theta}, p_{ref}) = -\mathbb{E}_{(x_1, x_2) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{p_{\theta}(x_1)}{p_{ref}(x_1)} - \beta \log \frac{p_{\theta}(x_2)}{p_{ref}(x_2)} \right) \right]$$

Easy for autoregressive

$$p(x_1, x_2, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

$$p(\text{The cat sat}) = p(\text{The})p(\text{cat}|\text{The})p(\text{sat}|\text{The cat})$$

Hard for diffusion

$$p(x_0) = \int p(x_{0:T}) dx_{1:T}$$

Which response do you prefer?
Your choice will help make ChatGPT better.

Response 1	Response 2
Correct. Explanation: The main categories of Microsoft 365 plans are Home, Business, and Enterprise. Next Question: Question 4: Which Microsoft 365 Enterprise plan includes advanced security and compliance features? • A. E1 • B. E2 • C. E3 • D. E5 What's your answer?	Correct. Explanation: The main categories of Microsoft 365 plans are Home, Business, and Enterprise. Next Question: Question 4: Which Microsoft 365 Enterprise plan includes advanced security and compliance features? • A. E1 • B. E2 • C. E3 • D. E5 What's your answer?


Hard to explore

Online: Reinforcement Learning

Agent interacts with an **environment** that provides a **reward**

Goal: Learn the policy the agent should follow that will maximize the agent's reward

$$p_{RL}^* = \underset{p}{argmax} \mathbb{E}_{x_{1:T} \sim p} \left[\sum_{i=1}^T R(x_i) \right]$$



Online: Reinforcement Learning

Train with a KL regularized version of the RL objective

$$\max_{\theta} \mathbb{E}_{x \sim p_{\theta}(x)} [R(x)] - \beta \cdot KL(p_{\theta} \parallel p_{ref})$$



Good reward



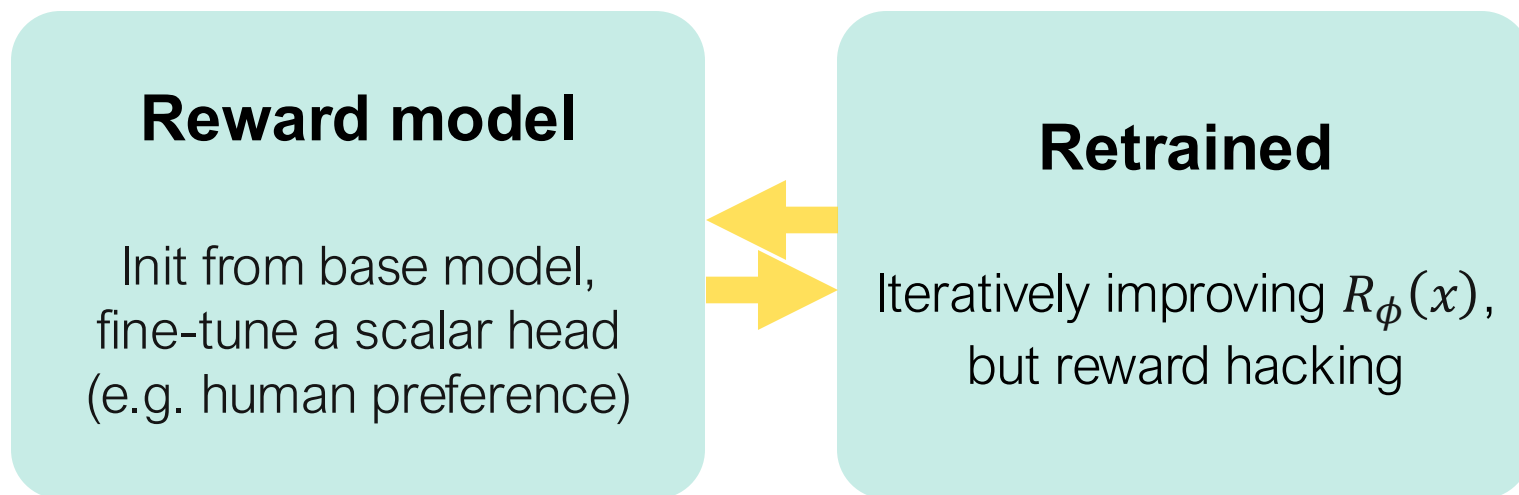
Don't go too far

The KL regularized RL objective converges to the reward tilted distribution!

$$p^*(x) = \frac{1}{Z} p_{ref}(x) \exp\left(\frac{R(x)}{\beta}\right)$$

Online: Getting reward the old way, RLHF

$$\min_{\phi} \mathbb{E}_{(x, y) \sim \mathcal{D}} \left[\|R_{\phi}(x) - y\|_2^2 \right]$$



Online: Getting reward with verifiers

$$R(x) = \begin{cases} 1 & \text{if } x \text{ is correct} \\ 0 & \text{if } x \text{ is wrong} \end{cases}$$

- No reward hacking
- Dense & scalable
- Rewards reasoning: Only correct final answers pay off, so need to discover multi-step strategies.



levent ✓
@__alpoge__

X.com

hello there the jacobian conjecture is false thanx to my close friend akhil for asking about it and my other close friend fable for working during the world cup final

$((1+xy)^3 z + y^2 (1+xy) (4+3xy), y + 3 x (1+xy)^2 z + 3 x y^2 (4+3xy), 2 x - 3 x^2 y - x^3 z):$
 $\setminus C^3 \setminus \text{to } \setminus C^3$, has jacobian determinant -2, and sends $(0, 0, -1/4)$, $(1, -3/2, 13/2)$, and $(-1, 3/2, 13/2)$ to $(-1/4, 0, 0)$

22:19 · 2026-07-19 · **20M** Views

1.3K 6.4K 32K 8.5K

Inference: Tilting at sampling time

Reward (or classifier-based) guidance: $\nabla \log p_0^*(x) = \nabla \log p_{ref,0}(x) + \frac{1}{\beta} \nabla R(x)$

Train a reward model on noisy versions of x , then add gradient to the unconditional score

Classifier-free guidance: $\nabla \log p_0(x)p_0(y|x)^w = (1 - w)\nabla p_{ref,0}(x) + w\nabla p_{ref,0}(x|y)$

Jointly train conditional & unconditional scores and interpolate between them during inference.

Does not actually converge to the true $p^*(x)$


Dhariwal, Prafulla, and Alexander Nichol. "Diffusion models beat gans on image synthesis." *NeurIPS* (2021).
Ho, Jonathan, and Tim Salimans. "Classifier-free diffusion guidance." *arXiv* (2022).

Inference: Fixing the bias with multiple samples

Feymann-Kac Steering

1. Sample N candidates
2. Drift at each step
3. Compute a weight w_i every step which says how good each candidate is
4. Resample the candidates based on w_i

This procedure converges to target $p^*(x)$ as $N \rightarrow \infty$!



*Does **not** require access to likelihoods, works under the hood by using Sequential Monte Carlo (SMC)*

Which one to pick?

Method	Cheap inference	Don't need to train	On or offline	Likelihoods help	Must add model capabilities	Other requirements
Prompt / condition / in-context	✓	✓	—	—	✗	—
SFT	✓	✗	Offline	✗	✗	Curated examples
Best-of-N / RAFT	✗	✓	Offline	✗	✗	Reward fxn
DPO	✓	✗	Offline	✓	✗	Preference pairs
RL: PPO / GRPO	✓	✗	Online	✓	✓	Scalar or verifiable reward
Guidance / FKC / Beam Search	✓ / ✗	✓	—	✗	✗	—

My 2 cents:


*Get a verifier & scale, or get a proxy reward & pray
Retrain when you can get a lot of data, otherwise, steer*

Outline

Part I. Simulation(-free) Generative Models

Part II. Fine-tuning & Steering

Part III. Applications to Chemistry



Number of Possible Proteins

$>10^{260}$






Atoms in the observable universe
 10^{80}

Evolution the tinkerer

Atoms in the observable universe
 10^{80}




● Proteins Nature Has Ever Tried
 $\sim 10^{44}$

Reading & Writing Nature's Language



Protein binders
(e.g. miniprotein, nanobodies)

What about nature's catalyst: enzymes?




More than nature:




The enzyme pill that makes you not lactose intolerant!

Or, degrade the PFAS in your water?


Directed evolution



Frances Arnold



The Problem




Hard
to find!


1. Natural enzymes excel in complicated chemistries, but they are made via evolutionary pressure.
2. Nature has only sampled a very narrow slice of all possible chemistry

Modern enzyme design: starting with a partial solution


1. Extract a known 'motif'



2. Generate protein backbone
with the motif fixed in space




3. Given a backbone structure,
predict a sequence




Why shouldn't we do this

1. No new motif = no new function
2. Sequence & structure should inform/constrain each other during design.
3. Steer our model based on sequence (cheaper data!)




So, what should we do?


1. Generate protein sequence & structure simultaneously
2. Build informed, new motifs on any biomolecule (ligands, DNA, RNA, TS)
3. Steer the model based on sequence and structure preferences



Building a multimodal model



Continuous diffusion



Discrete diffusion

We need to be able to sample from the joint marginals $p_{t,r}(x_t, y_r)$

Theory says  if we sample $t \sim U[0,1], r \sim U[0,1]$



Practice says      even if we do everything correct in theory!!!

Solving multimodal diffusion

Given a base model, how do we do inference?

Self-correction

Idea: AA prediction might be wrong early, we should let the model correct itself



Peng*, Bezemek*, ... "Path planning for masked diffusion model sampling." *arXiv (2025)*

Peng*, Bezemek*, ... "Planner Aware Path Learning in Diffusion Language Models Training." *ICLR Oral (2026)*


Solving multimodal diffusion

Given a base model, how do we do inference?

Entropy annealed tempering

Idea: Don't let the model be too confident early when not much structural detail, let it sharpen later

$$\tau_r^{(n)} = \tau \left(1 + \frac{r(H_{max} - H^{(n)})}{H_{max}} \right)$$




Solving multimodal diffusion

Given a base model, how do we do inference?



Noisy guidance

Idea: Condition the model on a noisier version of each modality to sharpen conditional predictions

$$\omega s_{t,r}^X(x_t, y_r) + (1 - \omega) s_{t,r+\Delta}^X(x_t, y_{r+\Delta})$$
$$\omega s_{t,r}^Y(x_t, y_r) + (1 - \omega) s_{t+\Delta,r}^Y(x_{t+\Delta}, y_r)$$



DISCO yields SOTA designs



Krishna, Rohith, et al. "Generalized biomolecular modeling and design with RoseTTAFold All-Atom." *Science* (2024)


Stark, Hannes, et al. "Boltzgen: Toward universal binder design." *bioRxiv* (2025)

Butcher, Jasper, et al. "De novo design of all-atom biomolecular interactions with rfdiffusion3." *bioRxiv* (2025)

Diverse and valid conformers!

FKC - Specificity Guidance

$t=1.000$



q^{on}




q^{off}




Generated

DISCO designs enzymes without fixed motifs




Only $\sim 10^4$ samples
1-2 orders magnitude fewer than SOTA


DISCO designs enzymes with new motifs




Designing motifs with chemical guidance



Zn metallohydrolase:
His/Asp/Glu coordinate to Zn



Retroaldase:
Lys N ζ $\rightarrow$ C13
& Tyr present



Morita-Baylis-Hillman:
His N ϵ 2 $\rightarrow$ C13

+ motif/scaffold novelty $\rightarrow$ separate section


Theophile 12:15 PM

We can call them discozymes that would sound good



New-to-nature carbene transferases

Alkene cyclopropanation



Hemin: 96 TTN

Evolved P450_{H2-5-F10}: 364 TTN


Heme-binding design C45: 455 TTN

Theozyme-based design PNC2: 630 TTN


Science **2013**, 339, 307

PNAS, **2020**, 1419

Science **2025**, 388, 665



C(*sp*³)-H alkylation




Hemin: ND


P-4 A82L (different substrate): 13 TTN

Evolved P411-CHF: 2030 TTN

Nature **2019**, 565, 67



B-H insertion




Hemin: 80 TTN


Rma cyt c: 120 TTN

Evolved *Rma* cyt c: 2490 TTN

Nature **2017**, 552, 132




Spirocyclopropanation



Hemin: ND

Evolved ApePgb: 2200 TTN

J. Am. Chem. Soc. **2025**, 147, 27165




Worse than evolution

Summary

- Training isn't everything!
- Solve multimodal diffusion by aligning the modalities at inference time: don't be overconfident.
- Design enzymes without a catalytic motif, and make it work in the lab!


Molecular crystal structure prediction

- Crystal packing directly governs the physical and chemical properties of organic solids.
- Goal: Given a *2D molecular graph*, can we predict its *3D crystal packing*?



Molecular crystal structure prediction

- Crystal packings occupy few local minima in a rugged Gibbs free *energy* landscape, where *kinetics* often dictate which minima are formed
- *Intramolecular* + *intermolecular* interactions.
- More diverse interactions than 20 amino acids




Classical way: via DFT

Step 1: Search



Enumeration, evolutionary algorithms, etc.

Step 2: Rank




Oracle access to energy models (force fields, DFT)

Limitations:


- Doesn't work on complicated samples
- Require *tens of thousands* of candidate structures per molecule, computationally prohibitively *expensive*

The challenges with CSP


1 *Joint distribution*
of thermodynamic
and kinetic effects



2 Diverse *long-range*
and *weak* molecular
interactions




3 Tough symmetry and
ambiguous Z/Z'



Aspirin
(Organic)


VS



Perovskite
(Inorganic)

OXtal blueprint

1 Joint distribution
of thermodynamic
and kinetic effects




Diffusion Model

2 Diverse long-range
and weak molecular
interactions

CCDC


~600K crystal structures

3 Tough symmetry +
ambiguous Z/Z'




Learn *local interactions* to
capture global symmetry

OXtal overview



Schematic of crystal growth




Weighted composite training objective:


$$\mathcal{L}(\theta) = \mathbb{E}_{t \sim \mathcal{U}(0,1), x_t \sim p_t(x_t | x_0^{\text{align}})} \left[\underbrace{\mathcal{L}_{\text{mse}}(\hat{x}_0, x_0^{\text{align}})}_{\text{Global accuracy}} + \underbrace{\mathcal{L}_{\text{sLDDT}}(\hat{x}_0, x_0^{\text{align}})}_{\text{Local environment}} \right] + \underbrace{\lambda_{\text{dist}} \mathcal{L}_{\text{dist}}(\hat{d}, d)}_{\text{Pairwise distances}}$$

OXtal generates accurate packings


QQQCIG
RMSD₁₅=0.55 Å






HURYUQ
RMSD₁₅=0.58 Å




TPHPOR
RMSD₁₅=1.23 Å




UMIMIO
RMSD₁₅=0.55 Å




GLYCIN
RMSD₁₅=0.74 Å




ACACPD
RMSD₁₅=1.13 Å



CAPRYL
RMSD₁₅=0.73 Å




QAXMEH
RMSD₁₅=1.24 Å




OXtal is competitive to DFT, but much faster

OXtal achieves competitive Pac_S while being *orders of magnitude cheaper* than DFT-based methods



Inference cost of OXtal vs DFT

DFT requires *significantly more samples*, suggesting that they are not sampling the correct energy wells



Sample efficiency (Target XXVIII in CSP7)

Summary

- OXtal is the *first* generative model to successfully scale to molecular crystals
- Reasons: diffusion model scales well when you just do data augmentation! Stable training is the key to success.



**generate
a valid
unit cell**



**generate
a locally
consistent,
periodic atomic
environment**

Overall summary

- Continuous-time generative models can now train stably in many domains.
- Particularly suited for chemical applications with a lot of reasonably quality data ($>10^4$) in a huge search space.
- Remaining challenges are:
 - scaling to LLM levels,
 - few step, especially in discrete space,
 - multimodal generation,
 - improving generalization to underlying distribution,
 - how to finetune with noisy rewards.